

CH 4: ADVANCED OPTIMIZATION ALGORITHMS FOR AI

PART I: ADAM AND ITS VARIANTS

Dong-Young Lim
UNIST

AI51101, IE55101


OBJECTIVE & ITS CHALLENGES

- ▶ **Objective:** Find a set of parameters θ that minimizes a high-dimensional and non-convex loss function $f(\theta)$.
- ▶ (Stochastic) Gradient Descent (GD) is a first-order iterative method that updates parameters based solely on the gradient at the current point (local information):

$$\theta_{t+1} = \theta_t - \gamma \nabla f(\theta_t)$$

- ▶ However, the loss landscapes of complex models present significant geometric challenges where this local approach is insufficient. Key issues include:
 - **Differential Curvature:** Different parameters have vastly different sensitivities.
 - **Flat Surfaces and Saddle Points:** Regions where the gradient is near-zero, severely slowing down convergence.
 - **Steep Cliffs:** Regions of excessively large gradients, leading to unstable updates.

CHALLENGE I: DIFFERENTIAL CURVATURE



- ▶ In many problems, the loss function forms a "ravine" - steep in one direction but gentle in another. This is called **differential curvature**.
- ▶ The left contour ($f = x^2 + y^2$) is a simple bowl. Any step downhill points to the center.
- ▶ The right contour ($f = x^2 + 5y^2$) is an elliptical bowl.
 - The gradient (steepest descent direction) points almost perpendicular to the ravine's floor.
 - This leads to a **zigzag pattern**: overshooting across the narrow axis while making slow progress along the main axis. A **single learning rate** can't be both large enough for the gentle slope and small enough for the steep slope.

CHALLENGE I: DIFFERENTIAL CURVATURE



(a) Basin



(b) Ravine


CHALLENGE I: DIFFERENTIAL CURVATURE

- ▶ **Question:** Can we solve the differential curvature problem by simply normalizing the input features?
- ▶ Normalization adjusts the scale of input features to be in a similar range. This can indeed help make the loss contours more uniform and is a recommended practice.
- ▶ **However, it is not a fundamental solution.**
 - The core problem of differential curvature is the varying sensitivity of the loss function with respect to each parameter θ .
 - Normalizing the inputs x is different from normalizing the parameter sensitivities θ . While the former can influence the latter, it does not directly solve the underlying issue.

CHALLENGE II: PLATEAUS & SADDLE POINTS


- ▶ **Saddle Points:** A point where the gradient is zero ($\nabla f(\theta) = 0$) but it is not a local extremum.
 - In high-dimensional functions, saddle points are exponentially more prevalent than local minima.
 - Standard GD algorithms are prone to get stuck near saddle points as the gradient diminishes.

- ▶ **Flat regions (Plateaus):** Nearly flat regions in the loss landscape.
 - In these regions, the gradient is consistently close to zero, which can halt the learning process. This is known as the **vanishing gradient problem**.




CHALLENGE III: STEEP CLIFFS

- ▶ The loss landscape can contain extremely steep regions, often called "cliffs".
- ▶ Descending a cliff can cause an excessively large update, leading to the **exploding gradient** problem and unstable training.



LOSS LANDSCAPE IN DEEP LEARNING



BEYOND STANDARD GRADIENT DESCENT

- ▶ As we have seen, the complex loss landscapes in deep learning present significant challenges for standard Gradient Descent.
- ▶ To overcome these challenges, various strategies that go **beyond the standard Gradient Descent** algorithm have been developed.
- ▶ We will explore two principal approaches:
 1. **Acceleration using Momentum**
 2. **Adaptive Learning Rates (ADAM and its variants)**

GRADIENT DESCENT WITH MOMENTUM

- ▶ **Key Idea:** Don't just rely on the local information, i.e., the current gradient. Accumulate past gradients to create "velocity" and accelerate the descent.
- ▶ **Physical Analogy:** A heavy ball rolling down a hill.
 - **Standard GD** is like a massless ball. It follows the steepest path at each point but stops instantly (No inertia).
 - **GD with Momentum** is like a rolling ball. It builds up speed (momentum) as it rolls, making it less affected by noisy gradients.
 - That's why GD with momentum is sometimes called "heavy-ball method"
- ▶ This accumulated velocity helps to:
 - Accelerate through flat plateaus where the gradient is small.
 - Dampen oscillations in narrow ravines.

GRADIENT DESCENT WITH MOMENTUM

- ▶ Momentum tracks an **exponential moving average (EMA)** of the gradients, which is treated as the current velocity.
- ▶ **Update Rules:** Let $g_t = \nabla f(\theta_t)$ (or, represent the stochastic gradient such that $\mathbb{E}[g_t] = \nabla f(\theta_t)$ where the randomness of the data is omitted.)
 1. Update velocity v_t using the previous velocity and current gradient:

$$v_t = \beta v_{t-1} + (1 - \beta)g_t$$


2. Update parameters θ_t :

$$\theta_{t+1} = \theta_t - \gamma v_t$$

- ▶ **Key Components:**
 - v_t : The "velocity" vector, representing the direction and speed of the update.
 - β : The momentum coefficient (e.g., 0.9). It controls how much of the past velocity is retained.
 - ▶ A high β means past gradients have a strong, lasting influence (less friction).
 - γ : The learning rate.

GRADIENT DESCENT WITH MOMENTUM

- ▶ **1. How it Navigates Ravines (Differential Curvature):**
 - Oscillating gradients (across the ravine) are in opposite directions and **cancel each other out** in the moving average.
 - Consistent gradients (along the ravine floor) point in the same direction and **accumulate**, building up velocity.
 - As a result, it reduces the zigzag motion and accelerates progress.
- ▶ **2. How it Escapes Plateaus & Saddle Points:**
 - When the optimizer encounters a flat region, the current gradient g_t becomes nearly zero.
 - However, the stored velocity v_t allows the optimizer to go through the flat area and push past the saddle point, where standard GD would stop.



FROM MOMENTUM TO ADAPTIVE RATES

- ▶ **Recap:** Momentum helps accelerate through flat regions and reduces oscillations.
- ▶ However, momentum still uses a **single & fixed learning rate** for all parameters.
- ▶ In a ravine, we ideally want to move slowly across the steep axis but quickly along the flat axis.
This requires different step sizes for different directions.
- ⇒ What if we could adapt the learning rate for each parameter individually? This leads us to our second strategy: **Preconditioning**.

FROM MOMENTUM TO ADAPTIVE RATES

- ▶ **Recap:** Momentum helps accelerate through flat regions and reduces oscillations.
 - ▶ However, momentum still uses a **single & fixed learning rate** for all parameters.
 - ▶ In a ravine, we ideally want to move slowly across the steep axis but quickly along the flat axis.
This requires different step sizes for different directions.
- ⇒ This brings us to our second principal strategy (adaptive learning rate):
- **Adaptive learning rate** scheme uses a different learning rate for each parameter at each iteration.
 - This is formally achieved through a mechanism known as **Preconditioning**.

PRECONDITIONING

- ▶ An optimization problem is **ill-conditioned** if the loss function's curvature dramatically varies across different directions.
- ▶ **Condition Number (κ):**
 - The degree of ill-conditioning is measured by the **condition number** of the Hessian:

$$\kappa(H) = \frac{\lambda_{\max}}{\lambda_{\min}}$$

where λ_{\max} and λ_{\min} are the largest and smallest eigenvalues of H .

- Eigenvalues represent curvature. A large κ means a steep and narrow ravine, which makes gradient descent less effective. On the other hand, a perfect circle has $\kappa = 1$.
- ▶ The goal of preconditioning is to transform an ill-conditioned problem into a well-conditioned one.

PRECONDITIONING

- ▶ **Key Idea:** Reshape the parameter space to make the loss contours more spherical (circular).
- ▶ This is achieved by multiplying the gradient by a matrix P_t^{-1} , called the **preconditioner**.
- ▶ **General Update Rule:**

$$\theta_{t+1} = \theta_t - \gamma P_t^{-1} g_t$$

The preconditioner P_t approximates the curvature of the loss function. Its inverse, P_t^{-1} , scales the gradient components: it shrinks steps in high-curvature directions and enlarges them in low-curvature directions.

- ▶ With preconditioning, parameter updates become balanced according to the local curvature of the loss surface.

PRECONDITIONING

- ▶ Consider the objective function:

$$f(x, y) = \frac{1}{2}(10x^2 + y^2)$$

with $\theta_0 = (1, 10)$.

- ▶ **1. Standard Gradient Descent Step**

- The gradient at this point is: $g_0 = \nabla f(1, 10) = [10x, y]|_{(1,10)} = [10, 10]$.

- ▶ **2. Preconditioned Step**

- Choose the Hessian as the preconditioner: $P = H = \begin{bmatrix} 10 & 0 \\ 0 & 1 \end{bmatrix}$.
- The preconditioned gradient is:

$$P^{-1}g_0 = \begin{bmatrix} 1/10 & 0 \\ 0 & 1 \end{bmatrix} \begin{bmatrix} 10 \\ 10 \end{bmatrix} = \begin{bmatrix} 1 \\ 10 \end{bmatrix}$$

- ▶ **Effect of Preconditioning:** preconditioning creates an adaptive learning rate- small for sensitive directions (high curvature) and large for stable directions (low curvature).

PRECONDITIONING

- ▶ In the previous toy example, we used the true Hessian matrix as the preconditioner. This represents the **ideal** case.
- ▶ **Ideal Preconditioner:** The Hessian (H_t)
 - The Hessian fully captures the curvature of the loss function, and using it ($P_t = H_t$) perfectly reshapes the problem space. This is known as **Newton's Method**.
 - **The Problem:** For a model with d parameters, computing and inverting the Hessian is computationally expensive and this is infeasible for large-scale neural networks.
- ▶ **Practical Approach: Diagonal Approximation**
 - In deep learning, we use cheap approximations. The most common is a **diagonal matrix**, which assumes zero correlation between parameters. This is the core idea behind:
 - ▶ AdaGrad, RMSProp, ADAM, and their variants

ADA GRAD

- ▶ AdaGrad (2011, JMLR) implements the diagonal preconditioning idea by the accumulated magnitude of past gradients.
- ▶ **Mechanism:**

- Let $g_t = \nabla f(\theta_t)$.

1. Accumulate squared gradients (element-wise):

$$s_t = s_{t-1} + g_t^2$$

2. Update parameters:

$$\theta_{t+1} = \theta_t - \frac{\gamma}{\sqrt{s_t + \epsilon}} g_t$$

- ▶ It performs well on problems with sparse features, as infrequent parameters receive larger updates.
- ▶ **Limitation:** The accumulator s_t grows monotonically. As a result, the effective learning rate for all parameters steadily decreases and eventually becomes very small, causing the model to stop learning prematurely.

RMSPROP

- ▶ RMSProp (2012, unpublished) addresses AdaGrad's critical weakness by preventing the learning rate from monotonically decreasing.
- ▶ **Mechanism:**
 - Instead of summing all past squared gradients, it uses an **exponential moving average (EMA)**, which allows old gradients to be "forgotten".
 - 1. Update the EMA of squared gradients (element-wise):

$$s_t = \rho s_{t-1} + (1 - \rho) g_t^2, \quad s_{-1} = 0$$

- 2. Update parameters:

$$\theta_{t+1} = \theta_t - \frac{\gamma}{\sqrt{s_t + \epsilon}} g_t$$

- ▶ Here, ρ is a decay rate, typically set to 0.9 or 0.99. By keeping a moving average, the denominator does not grow indefinitely.
- ▶ **Limitation:** The accumulator s_t is initialized to zero, making it biased towards zero in the early stages of training. This can cause unintentionally large parameter updates initially.

ADAM

- ▶ Adam (2015, ICLR) is arguably the most popular adaptive optimizer, effectively combining the ideas of **Momentum** and **RMSProp**.
- ▶ It computes and maintains two separate exponential moving averages:
 1. **First Moment** (m_t): EMA of the gradients like Momentum.
 2. **Second Moment** (v_t): EMA of the squared gradients (like RMSProp).
- ▶ **The Full Adam Algorithm:**
 1. Compute gradient: $g_t = \nabla f(\theta_{t-1})$
 2. Update 1st moment: $m_t = \beta_1 m_{t-1} + (1 - \beta_1)g_t$
 3. Update 2nd moment: $v_t = \beta_2 v_{t-1} + (1 - \beta_2)g_t^2$
 4. Correct bias in 1st moment: $\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$
 5. Correct bias in 2nd moment: $\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$
 6. Update parameters: $\theta_t = \theta_{t-1} - \gamma \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$

ADAM PROPERTY I: BIAS CORRECTION

- ▶ **The Problem:** The moment vectors, m_t and v_t , are initialized as zeros. This causes them to be biased towards zero, especially during the initial steps of training.
- ▶ **The Solution:** Adam introduces bias-correction terms that analytically compute the expected value of the moments and scale them appropriately.

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t} \quad , \quad \hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- ▶ By calculating the expectation, we see that

$$\mathbb{E}[\hat{m}_t] = \mathbb{E}[g_t]$$

and

$$\mathbb{E}[\hat{v}_t] = \mathbb{E}[g_t^2]$$

- ▶ This correction ensures that the update step has a proper magnitude from the very beginning of training.

ADAM PROPERTY I: BIAS CORRECTION

- ▶ Assume a stationary process with $\mathbb{E}[g_t] = \mu$ and $\mathbb{E}[g_t^2] = \nu$ in the early phase.

$$\mathbb{E}[m_t] = (1 - \beta_1) \sum_{k=1}^t \beta_1^{t-k} \mu = (1 - \beta_1^t) \mu \Rightarrow \mathbb{E}[\hat{m}_t] = \mu.$$

$$\mathbb{E}[v_t] = (1 - \beta_2) \sum_{k=1}^t \beta_2^{t-k} \nu = (1 - \beta_2^t) \nu \Rightarrow \mathbb{E}[\hat{v}_t] = \nu.$$

- ▶ Thus \hat{m}_t and \hat{v}_t are nearly unbiased in the early phase of training.

ADAM PROPERTY II: SCALE INVARIANCE

- ▶ The magnitude of Adam's update step is largely invariant to the scale of the gradients.
- ▶ If we rescale the gradient by a constant factor c (i.e., $g'_t = c \cdot g_t$), how does the update change?
 - The first moment estimate scales by c : $\hat{m}'_t = c \cdot \hat{m}_t$
 - The second moment estimate scales by c^2 : $\hat{v}'_t = c^2 \cdot \hat{v}_t$
- ▶ The final update term becomes:

$$\frac{\hat{m}'_t}{\sqrt{\hat{v}'_t + \epsilon}} \approx \frac{c \cdot \hat{m}_t}{\sqrt{c^2 \cdot \hat{v}_t}} = \frac{c \cdot \hat{m}_t}{|c| \cdot \sqrt{\hat{v}_t}} = \text{sign}(c) \frac{\hat{m}_t}{\sqrt{\hat{v}_t}}$$

- ▶ **Implication:** The magnitude of the update is unaffected by the scaling factor c . This makes Adam **robust** to objective functions that produce very large or very small gradients.

ADAM PROPERTY III: BOUNDED STEP SIZE

- ▶ Each coordinate's step is effectively capped by γ in stationary regimes.
- ▶ The per-coordinate update is:

$$|\Delta\theta_{t,i}| = \gamma \frac{|\hat{m}_{t,i}|}{\sqrt{\hat{v}_{t,i}} + \epsilon}$$

By Cauchy–Schwarz on EMAs (stationary regime),

$$|\hat{m}_{t,i}| \leq \sqrt{\hat{v}_{t,i}} \quad \Rightarrow \quad |\Delta\theta_{t,i}| \leq \gamma \frac{\sqrt{\hat{v}_{t,i}}}{\sqrt{\hat{v}_{t,i}} + \epsilon} \leq \gamma.$$

- ▶ **Intuition:** Adam has a built-in trust region property. It prevents the optimizer from taking excessively large steps, which greatly improves training stability, especially in regions with noisy or large gradients.

ADAM PROPERTY IV: NOISE-TO-SIGNAL RATIO

- ▶ We can gain a deeper intuition for Adam's behavior by viewing it through the lens of the Noise-to-Signal Ratio (NSR):
 - **Signal:** The consistent direction of descent, estimated by the mean of the gradients ($m_t \approx \mathbb{E}[g_t]$).
 - **Noise:** The uncertainty or fluctuation of the gradients, measured by their standard deviation ($\sqrt{\text{Var}(g_t)}$).
 - NSR measures the relative uncertainty of the gradient for a given parameter:
- NSR =
$$\frac{\sqrt{\text{Var}(g_t)}}{|\mathbb{E}[g_t]|}$$
- A **high NSR** indicates a noisy, unreliable gradient, while a **low NSR** indicates a consistent, reliable gradient.
- ▶ How Adam Responds to NSR:
 - When **NSR is high** (high noise), the variance term dominates \hat{v}_t . This increases the denominator, which in turn **reduces the step size**.
 - When **NSR is low** (high signal), the mean term dominates, leading to a relatively **larger step size**.

ADAM: DEFAULT SETTING

► Adam:

1. Compute gradient: $g_t = \nabla f(\theta_{t-1})$
2. Update 1st moment: $m_t = \beta_1 m_{t-1} + (1 - \beta_1)g_t$
3. Update 2nd moment: $v_t = \beta_2 v_{t-1} + (1 - \beta_2)g_t^2$
4. Correct bias in 1st moment: $\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$
5. Correct bias in 2nd moment: $\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$
6. Update parameters: $\theta_t = \theta_{t-1} - \gamma \frac{\hat{m}_t}{\sqrt{\hat{v}_t + \epsilon}}$

► Default Setting:

$$\gamma = 10^{-3}, \beta_1 = 0.9, \beta_2 = 0.999, \epsilon = 10^{-8}.$$

ADAM: THE CONVERGENCE

- ▶ The original Adam paper provided a proof of convergence for online convex optimization, showing a favorable $O(\sqrt{T})$ regret bound.
- ▶ However, subsequent research identified a flaw in the proof. It was demonstrated through counterexamples that **Adam can fail to converge** even on simple convex problems in the paper of AMSGrad.
- ▶ Since this flaw, analyzing the (non-)convergence of Adam has become a interesting topic of research in the optimization community.

On the convergence of adam and beyond

SJ Reddi, S Kale, S Kumar - arXiv preprint arXiv:1904.09237, 2019 - arxiv.org

... only prove **non-convergence** of ADAM in the ... of ADAM and AMSGRAD on synthetic example on a simple one dimensional convex problem inspired by our examples of **non-convergence**...

☆ 저장 99 인용 3666회 인용 관련 학술자료 전체 9개의 버전 ☰

Non-convergence and limit cycles in the adam optimizer

S Bock, M Weiß - International Conference on Artificial Neural Networks, 2019 - Springer

... This is done for the **Adam** algorithm in batch mode only, but ... to clarify the global **non-convergence** of **Adam** even for strictly ... 2 we define our variant of the **Adam** algorithm and explain ...

☆ 저장 99 인용 18회 인용 관련 학술자료 전체 7개의 버전 ☰

Non-convergence of Adam and other adaptive stochastic gradient descent optimization methods for non-vanishing learning rates

S Dereich, R Graeber, A Jentzen - arXiv preprint arXiv:2407.08100, 2024 - arxiv.org

... and the **Adam** ... **Adam** optimizer fail to converge to any possible random limit point if the learning rates are asymptotically bounded away from zero. In our proof of this **non-convergence** ...

☆ 저장 99 인용 12회 인용 관련 학술자료 전체 4개의 버전 ☰

On the convergence of a class of adam-type algorithms for non-convex optimization

AMSGRAD

- ▶ We will examine a simple 1D **convex** problem where Adam is shown to converge to a suboptimal solution.
- ▶ Problem Setup: Online Convex Optimization
 - **Domain:** $\mathcal{F} = [-1, 1]$.
 - **Function Sequence:** A repeating 3-step sequence of functions.

$$f_t(x) = \begin{cases} Cx & \text{if } t \mod 3 = 1 \\ -x & \text{otherwise} \end{cases} \quad (\text{for a large constant } C > 2)$$

- **Gradients:** g_t is C once, then a small negative value -1 twice.
- ▶ **Goal:** Find the point that minimizes the cumulative loss:

$$F_{cycle}(x) = Cx + (-x) + (-x) = (C - 2)x$$

Since $C > 2$, the term $(C - 2)$ is positive. To minimize this on the domain $[-1, 1]$, we must choose the smallest possible value -1 for x . (You need to study online learning for full details.)

DETOUR: A REGRET ANALYSIS

- ▶ **The Goal of Online Learning:** The algorithm doesn't know the future loss functions. Its goal is to make a sequence of choices (x_1, x_2, \dots) that, in total, performs nearly as well as a hypothetical "expert" who knew everything in advance.
- ▶ We define this expert's performance as the best possible outcome using a single and fixed point x^* over the entire history. This provides a stable and consistent **benchmark** to measure our algorithm against. The ultimate goal of training a model is often to find one final set of parameters, so this benchmark is very practical.
- ▶ We measure performance by **Regret**, R_T , which is the difference between our algorithm's total loss and the expert's total loss.

$$R_T = \underbrace{\sum_{t=1}^T f_t(x_t)}_{\text{Algorithm's Total Loss}} - \underbrace{\min_{x^* \in \mathcal{F}} \sum_{t=1}^T f_t(x^*)}_{\text{Best Fixed Point's Total Loss}}$$

- ▶ A good algorithm ensures that its average regret, R_T/T , approaches zero as $T \rightarrow \infty$. This means our algorithm, on average, performs as well as the best fixed point in hindsight.

AMSGRAD

- ▶ To demonstrate the failure, we use specific parameters that simplify the analysis: $\beta_1 = 0$ (so $m_t = g_t$) and a small β_2 .
- ▶ Case 1: The Normal Step ($g_t = -1$)
 - Most of the time, the gradient is -1. The second moment estimate v_t (the EMA of squared gradients) will therefore converge to $(-1)^2 = 1$ when β_2 is very small.
 - The parameter update is:

$$\Delta\theta = -\gamma \frac{g_t}{\sqrt{v_t + \epsilon}} \approx -\gamma \frac{-1}{\sqrt{1}} = +\gamma$$

- **Observation:** During these iterations, small-gradient steps, Adam moves towards +1.

AMSGRAD

► Case 2: The "Spike" Step ($g_t = C$)

- When the rare, large gradient C appears, v_t spikes.
- For a small β_2 , the denominator $\sqrt{v_t}$ becomes approximately C .
- The parameter update is:

$$\Delta\theta = -\gamma \frac{g_t}{\sqrt{v_t + \epsilon}} \approx -\gamma \frac{C}{C} = -\gamma$$

- **Observation:** The large and informative gradient is scaled down so much that it only produces a small step of size γ in the correct direction.

AMSGRAD

- ▶ Then, the net change over one 3-step cycle:
 - Step 1 (Spike): $\Delta\theta \approx -\gamma$
 - Step 2 (Normal): $\Delta\theta \approx +\gamma$
 - Step 3 (Normal): $\Delta\theta \approx +\gamma$
- ▶ Total Change over 3 steps $\approx (-\gamma) + (+\gamma) + (+\gamma) = +\gamma$
- ▶ Adam consistently drifts in the wrong direction (towards +1), away from the true minimum at -1.
- ▶ This shows a fundamental flaw of ADAM under a rare and large gradients environment.

AMSGRAD

- ▶ AMSGrad (2018) fixes this by adding a "**long-term memory**" to the second moment, guaranteeing a non-increasing adaptive learning rate.
- ▶ **The Key Modification:**



- AMSGrad maintains the maximum of the second moment estimates seen so far:

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\hat{v}_t = \max(\hat{v}_{t-1}, v_t)$$

- ▶ **How this solves the non-convergence issue:**
 - By using the running maximum, AMSGrad ensures the denominator used for scaling never shrinks.
 - Once it encounters a large squared gradient (like C^2 from our example), the denominator will remain large for all future steps. This stabilizes the adaptive learning rate, preventing it from becoming inappropriately large.

AMSGRAD



AMSGRAD

- ▶ There exist convex problems where Adam has non-zero average regret under $\beta_1 < \sqrt{\beta_2}$.

Theorem 2. For any constant $\beta_1, \beta_2 \in [0, 1)$ such that $\beta_1 < \sqrt{\beta_2}$, there is an online convex optimization problem where ADAM has non-zero average regret i.e., $R_T/T \not\rightarrow 0$ as $T \rightarrow \infty$.

- ▶ AMSGrad achieves $R(T) = O(\sqrt{T})$ in convex settings.

Theorem 4. Let $\{x_t\}$ and $\{v_t\}$ be the sequences obtained from Algorithm 2, $\alpha_t = \alpha/\sqrt{t}$, $\beta_1 = \beta_{11}$, $\beta_{1t} \leq \beta_1$ for all $t \in [T]$ and $\gamma = \beta_1/\sqrt{\beta_2} < 1$. Assume that \mathcal{F} has bounded diameter D_∞ and $\|\nabla f_t(x)\|_\infty \leq G_\infty$ for all $t \in [T]$ and $x \in \mathcal{F}$. For x_t generated using the AMSGARD (Algorithm 2), we have the following bound on the regret

$$R_T \leq \frac{D_\infty^2 \sqrt{T}}{\alpha(1 - \beta_1)} \sum_{i=1}^d \hat{v}_{T,i}^{1/2} + \frac{D_\infty^2}{(1 - \beta_1)^2} \sum_{t=1}^T \sum_{i=1}^d \frac{\beta_{1t} \hat{v}_{t,i}^{1/2}}{\alpha_t} + \frac{\alpha \sqrt{1 + \log T}}{(1 - \beta_1)^2(1 - \gamma)\sqrt{(1 - \beta_2)}} \sum_{i=1}^d \|g_{1:T,i}\|_2.$$

ADAMW

- ▶ In standard SGD, there are two common ways to implement regularization:
 - **L_2 Regularization:** Add the L_2 penalty term to the objective.

$$\min_{\theta} f(\theta) + \lambda \|\theta\|_2^2$$

Therefore, the gradient and the update rule become

$$g_t = \nabla f(\theta_t) + \lambda \theta_t$$

$$\theta_{t+1} = \theta_t - \gamma g_t$$

- **Weight Decay:** Directly subtract a fraction of the weights from the parameters after the gradient step.

$$\theta_{t+1} = (1 - \gamma \lambda) \theta_t - \gamma \nabla f(\theta_t)$$

- ▶ **For SGD, these two forms are identical.** This equivalence has led many to use the terms interchangeably. However, we will see this is not true for adaptive optimizers.

ADAMW

```
optimizer = torch.optim.Adam(lr=0.001, weight_decay=1e-5)
```

ADAMW

- ▶ Adam uses a diagonal preconditioner, $D_t = \text{diag}((\hat{v}_t + \epsilon)^{-1/2})$, to scale the gradients.
- ▶ **Naive L_2 -regularization to ADAM:** When L_2 regularization is naively added to the gradient *before* preconditioning, the update becomes:

$$\theta_{t+1} = \theta_t - \gamma D_t (\hat{m}_t + \lambda \theta_t)$$

- ▶ **Critical Flaw:** The regularization term $\lambda \theta_t$ is now also scaled by the adaptive matrix D_t .
 - This is no longer equivalent to minimizing a standard L_2 penalty, $\frac{\lambda}{2} \|\theta\|_2^2$.
 - Instead, it's equivalent to minimizing a complex, time-varying weighted penalty: $\frac{\lambda}{2} \theta^\top D_t \theta$.
- ▶ Parameters with large historical gradients will have a large \hat{v}_t , making their corresponding entry in D_t small. As a result, they receive **less regularization**, which is the opposite of the desired regularization effect.

ADAMW

- ▶ AdamW solves this by **decoupling** the weight decay from the adaptive gradient update.
- ▶ **AdamW Update Rule:**

$$\begin{aligned}\theta_{t+1} &= (1 - \gamma\lambda)\theta_t - \gamma D_t \hat{m}_t \\ &= \theta_t - \gamma D_t \hat{m}_t - \gamma\lambda\theta_t.\end{aligned}$$

1. First, the adaptive gradient step is calculated based on the moments: $-\gamma D_t \hat{m}_t$.
 2. Separately, the weight decay is applied uniformly to all weights: $-\gamma\lambda\theta_t$.
- ▶ This restores the intended effect of weight decay. The amount of weight shrinkage is now independent of the gradient history, leading to more stable and predictable regularization.
 - ▶ Default Settings: $\gamma = 0.001$, $\lambda = 0.01$.

ADAMW

- ▶ **Proximal view:** The AdamW update rule can be derived as the solution to the following optimization problem at each step:

$$\theta_{t+1} = \arg \min_{\theta} \underbrace{\langle \hat{m}_t, \theta - \theta_t \rangle + \frac{1}{2\gamma} \|\theta - \theta_t\|_{D_t^{-1}}^2}_{\text{Adaptive Gradient Step}} - \underbrace{\gamma \lambda \theta_t}_{\text{L2 Penalty}}$$

where $\|\theta\|_{D_t^{-1}} := \theta^\top D_t^{-1} \theta$.

- The first two terms aim to follow the momentum (\hat{m}_t) but penalize moving too far from the current point (θ_t).
 - The third term is the standard L_2 penalty.
- ▶ Taking the first-order condition of this objective and solving for θ_{t+1} gives us the exact AdamW update rule:

$$\theta_{t+1} = \theta_t - \gamma D_t \hat{m}_t - \gamma \lambda \theta_t$$


A NEW QUESTION: DO FASTER OPTIMIZERS GENERALIZE BETTER?

- ▶ By the late 2010s, Adam and its variants became the mainstream choice for training deep neural networks due to their remarkably fast training speed.
- ▶ **Emerging Problems:**
 - However, models trained with Adam, despite reaching low training error quickly, often showed worse performance on the **test set** compared to models trained with simple, slower SGD.
 - The "best" optimizer for training loss was not always the "best" for test accuracy.
- ▶ **Critical Questions:**
 - Does the choice of optimizer influence not just *how fast* we find a solution, but *which* solution we find?
 - And do different optimizers converge to solutions with fundamentally different **generalization abilities**?

SGD vs ADAM-LIKE OPTIMIZERS

- ▶ A seminal paper by Wilson et al. (2017) provided a striking answer to this question.
- ▶ **Main Claim:** Adaptive methods like Adam can find fundamentally different solutions from SGD, and these solutions often **generalize significantly worse**.
- ▶ **Key Findings:**
 - They constructed a simple convex problem where SGD achieves perfect generalization (0% test error), while Adam fails completely (50% test error).
 - They showed empirically across multiple deep learning tasks that well-tuned SGD outperforms well-tuned Adam on test performance, even when Adam's training loss is lower.


SGD vs ADAM-LIKE OPTIMIZERS




(a) CIFAR-10 (Train)

(b) CIFAR-10 (Test)

SGD vs ADAM-LIKE OPTIMIZERS




(a) War and Peace (Training Set)




(b) War and Peace (Test Set)

SGD vs ADAM-LIKE OPTIMIZERS




(c) Discriminative Parsing (Training Set)




(d) Discriminative Parsing (Development Set)

SGD vs ADAM-LIKE OPTIMIZERS



(e) Generative Parsing (Training Set)



(f) Generative Parsing (Development Set)

SGD vs ADAM-LIKE OPTIMIZERS

- ▶ **Why?** In overparameterized models, there are many possible solutions that achieve zero training error.
- ▶ The optimization algorithm's design gives it an "**implicit bias**" that makes it prefer certain types of solutions over others:
 - **SGD's Implicit Bias:**
 - ▶ When started from zero, SGD has an implicit bias towards finding the solution with the **minimum Euclidean (L_2) norm**.
 - ▶ This minimum norm solution is often associated with large margins and good generalization.
 - **Adaptive Methods' Implicit Bias:**
 - ▶ Their preconditioner gives them a different bias. They don't converge to the min L_2 norm solution.
 - ▶ They sometimes tend to converge to solutions that rely on non-generalizable features.

SGD vs ADAM-LIKE OPTIMIZERS

- ▶ The concept of "**implicit bias**" provides one compelling explanation for the generalization gap observed between SGD and adaptive methods.
- ▶ **However, this is just one possible explanation.**
- ▶ The deep and complex relationship between the choice of optimizer and a model's final generalization ability is still not fully understood. It remains a somewhat **mysterious and open problem**.
- ▶ Understanding exactly why different algorithms lead to solutions with different test performances is one of the most significant and active areas of research in deep learning theory today.