

CH 5: ADVANCED OPTIMIZATION ALGORITHMS FOR AI

PART II: SHARPNESS-AWARE MINIMIZATION (SAM)

Dong-Young Lim
UNIST

AI51101, IE55101

MOTIVATION: BEYOND ADAM

- ▶ Adam is a very popular and powerful optimizer because it is fast and stable. For a long time, it was the standard choice for training deep networks.
- ▶ However, researchers started to notice a common problem. Models trained with Adam often had **worse test performance** than models trained with simple SGD.
- ▶ This raised an important question for the field of optimization:

Why do some optimizers generalize better than others?

WHAT AFFECTS GENERALIZATION?

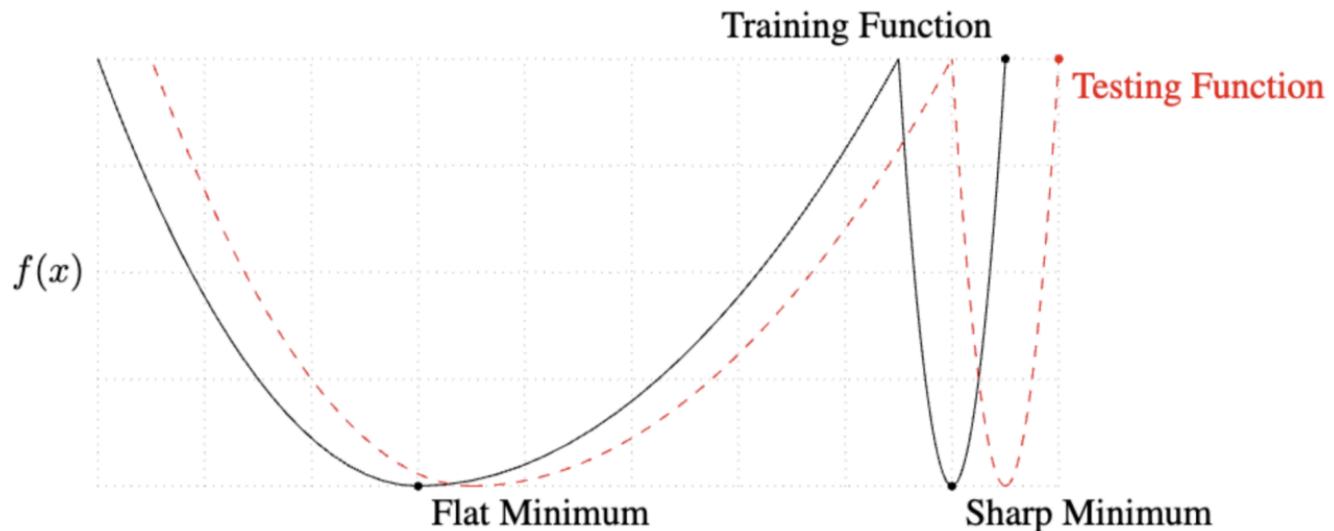
- ▶ Empirical studies have found many factors that influence generalization:
 - Model capacity and overparameterization
 - Regularization and noise injection (like Dropout)
 - Batch size and learning rate schedules
 - The "implicit bias" of the optimization algorithm
- ▶ Among them, one concept seemed to connect many of these observations:

Sharpness of the loss landscape.
- ▶ Intuitively, sharpness measures how "fragile" a solution is. A solution in a sharp minimum is very sensitive to small changes in the parameters, while a solution in a flat minimum is robust.

AN OLD IDEA REVISITED: FLAT MINIMA

- ▶ The idea that flat minima lead to better generalization is not new. It was proposed as early as 1997 in the paper "Flat Minima" (Hochreiter & Schmidhuber, 1997).
- ▶ **Main Idea:**
 - They argued that the training set is just a small sample of the true data distribution. The loss landscape of the training set is therefore a noisy version of the true loss landscape.
 - A sharp minimum on the training set might not be a minimum at all on the true landscape.
 - A flat minimum, however, represents a large region of low error, making it much more likely to be a low-error region on the true landscape as well.

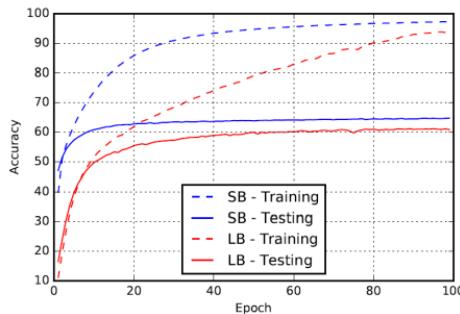
AN OLD IDEA REVISITED: FLAT MINIMA



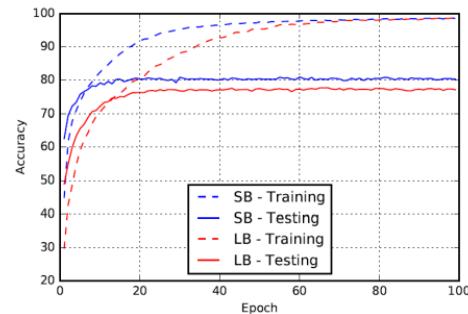
THE SHARPNESS HYPOTHESIS IN DEEP LEARNING

- ▶ The generalization gap between Adam and SGD made researchers look again at the old "flat minima" idea from 1997.
- ▶ The "Sharpness Hypothesis" returned as the main explanation for this gap.
- ▶ The hypothesis states that the geometry of a minimum is the key to generalization in modern deep networks.
- ▶ Starting from 2017, a series of important papers provided the strong empirical evidence for this hypothesis. To understand this, we will review three key papers in order:
 1. (2017 ICLR) On Large-Batch Training: Generalization Gap and Sharp Minima
 2. (2017 NeurIPS) Train Longer, Generalize Better
 3. (2020 ICLR) Fantastic Generalization Measures and Where to Find Them

- ▶ **Core Observation:** Training with a large batch size (e.g., > 1024) leads to a significant drop in test accuracy compared to small-batch training (e.g., 256).
- ▶ This "generalization gap" happened even when both methods reached the same low training error.
- ▶ Importantly, this was not standard over-fitting. The test accuracy for large-batch models would simply plateau at a lower value, not decrease after peaking.



(a) Network F_2



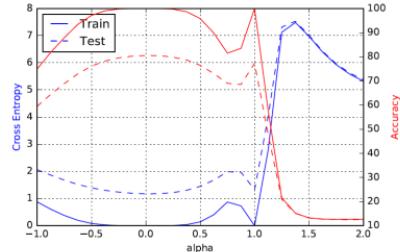
(b) Network C_1

- ▶ To find the cause, the authors investigated the **geometry** of the solutions.
- ▶ **Linear Interpolation**
 - First, find a solution from small-batch training, θ_{SB}^* , and one from large-batch training, θ_{LB}^* .
 - Then, define a line in the parameter space that connects these two points:

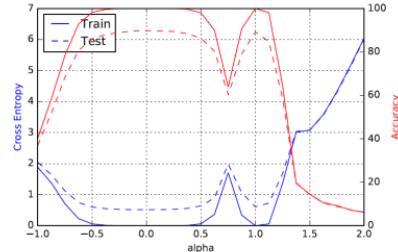
$$\theta(\alpha) = (1 - \alpha)\theta_{SB}^* + \alpha\theta_{LB}^*$$

- They plotted the loss function $f(\theta(\alpha))$ for values of α in a range like $[-1, 2]$. This shows the landscape along the line connecting the two solutions and beyond.

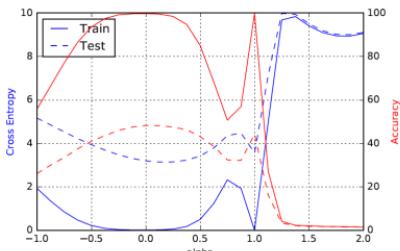
(2017 ICLR) KESKAR ET AL.



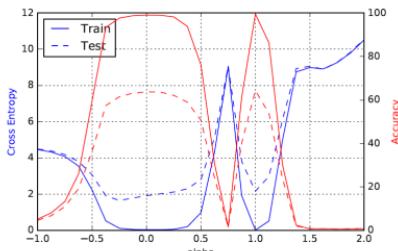
(c) C_1



(d) C_2



(e) C_3



(f) C_4

- ▶ The plots consistently showed that the LB solution ($\alpha = 1$) was at the bottom of a much **sharper**, narrower valley, while the SB solution ($\alpha = 0$) was in a wider, flatter basin.

(2017 ICLR) KESKAR ET AL.

- ▶ To support the visualizations, they proposed a metric to quantify the sharpness of a minimum.
- ▶ The formula is defined as:

$$\text{Sharpness}(\theta^*) = \frac{\max_{\|\epsilon\| \leq \rho} (L(\theta^* + \epsilon) - L(\theta^*))}{1 + L(\theta^*)}$$

The metric measures the highest loss value the function can reach within a small neighborhood around a given solution θ^* .

(2017 ICLR) KESKAR ET AL.

	$\epsilon = 10^{-3}$		$\epsilon = 5 \cdot 10^{-4}$	
	SB	LB	SB	LB
F_1	1.23 ± 0.83	205.14 ± 69.52	0.61 ± 0.27	42.90 ± 17.14
F_2	1.39 ± 0.02	310.64 ± 38.46	0.90 ± 0.05	93.15 ± 6.81
C_1	28.58 ± 3.13	707.23 ± 43.04	7.08 ± 0.88	227.31 ± 23.23
C_2	8.68 ± 1.32	925.32 ± 38.29	2.07 ± 0.86	175.31 ± 18.28
C_3	29.85 ± 5.98	258.75 ± 8.96	8.56 ± 0.99	105.11 ± 13.22
C_4	12.83 ± 3.84	421.84 ± 36.97	4.07 ± 0.87	109.35 ± 16.57

- ▶ For all tested networks, the sharpness value for the LB solutions was **10 to 100 times higher** than for the SB solutions.

- ▶ The high gradient noise in small-batch SGD prevents the optimizer from getting stuck in narrow (sharp) valleys.
- ▶ The low noise in large-batch SGD allows it to easily converge to the nearest minimum, which is often sharp.

In short, the randomness in SB-SGD helps it explore the landscape and find better, flatter solutions.

(2017 NEURIPS) HOFFER ET AL.

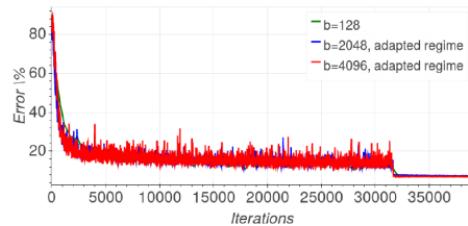
- ▶ This paper directly challenged the conclusions of Keskar et al.
- ▶ **Core Question:** Is the generalization gap really caused by the batch size itself, or is the comparison between SB and LB training unfair?
- ▶ **Hypothesis:** The key difference is not the batch size, but the total number of parameter updates. An epoch of SB training involves many more updates than an epoch of LB training.
- ▶ **Main Message:** The poor performance of large-batch models is simply a result of under-training.

- ▶ The authors designed experiments to give LB training a fair chance by matching the total training budget.
- ▶ **1. "Adapted Regime":**
 - Their main idea was to ensure the same number of gradient steps for all models.
 - If a large batch is k times bigger than a small batch (e.g., $B_{LB} = k \cdot B_{SB}$), they trained the large-batch model for k times as many epochs.
- ▶ **2. Adjusting Hyperparameters:**
 - They also argued that to keep the optimization dynamics similar, the learning rate should be scaled up for larger batches.

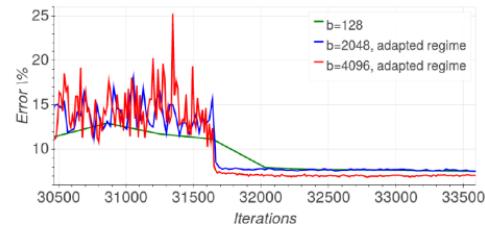
$$\eta_{LB} = \eta_{SB} \times \sqrt{k}$$

.

- ▶ **Main Findings:** When the total number of updates and hyperparameters were properly matched, the generalization gap between small-batch and large-batch training **completely disappeared**.
- ▶ Large-batch methods were able to achieve the same (or even slightly better) test accuracy as small-batch methods.
- ▶ This strongly suggested that there is no fundamental problem with large batches.



(a) Validation error



(b) Validation error - zoomed

(2017 NEURIPS) HOFFER ET AL.

- ▶ **Conclusion:** The generalization gap is not inherent to batch size. The "bad generalization of large batch" is due to the optimization schedule.
- ▶ This work reframed the debate. Sharpness is not a fundamental *cause* of bad generalization.
- ▶ This is a classic "correlation vs. causation" problem that is very difficult to solve completely.
- ▶ This debate created a clear need for a larger, more definitive study to test which properties are truly predictive of generalization, regardless of the training setup.

(2020 ICLR) JIANG ET AL.

- ▶ After the debate (NeurIPS 2017, Hoffer et al), this paper aimed to provide a data-driven answer.
- ▶ **Goal:** To test which of the many theories about generalization (norms, margins, sharpness, etc.) actually work in practice.
- ▶ **Methodology:**
 - They trained over 10,000 models by changing many different hyperparameters (depth, width, optimizer, etc.).
 - All models were trained to the exact same final training loss for a very fair comparison.
 - They then checked which of 40+ "generalization measures" best correlated with the final test performance.

► **Finding 1:** The Failure of Norms

- Measures based on the size of the weights (like L_2 norm or spectral norm) were very poor predictors of generalization.
- Sometimes, they were even negatively correlated, meaning larger weights led to better generalization. This contradicts many classical theories.

► **Finding 2:** The Success of Sharpness

- One category of measures consistently performed the best.
- Measures related to **sharpness** were the most reliable and accurate predictors of the generalization gap across all 10,000+ experiments.

- ▶ **Conclusion:** This paper provided strong empirical evidence that the sharpness of a minimum is a fundamental property linked to generalization.
- ▶ Even when controlling for training time and final loss, sharpness remains the best predictor of how well a model will perform.
- ▶ It is not just a side effect of training; it is a reliable indicator of a good solution.
- ▶ **Implication:** This work established that finding flat minima is a valid and important goal for developing better optimization algorithms.

SAM: MOTIVATION

- ▶ The previous research (especially ICLR 2020, Jiang et al.) gave us a clear goal: **we should find flat minima.**
- ▶ Standard optimizers like SGD or Adam do not directly search for flatness. They only try to minimize the loss at the current point w :

$$\min_w L_S(w)$$

- ▶ **Main Idea of SAM:** Instead of just minimizing the loss at a single point, let's change the optimization objective to find a point w that has low loss over an entire **neighborhood**. This will force the optimizer to find flat regions.
- ▶ Note that we will follow the notations used in the corresponding papers.

SAM: THE OBJECTIVE

- ▶ SAM changes the standard loss function $L_S(w)$ to a sharpness-aware loss function, $L_S^{\text{SAM}}(w)$:

$$\min_w L_S^{\text{SAM}}(w) \quad \text{where} \quad L_S^{\text{SAM}}(w) = \max_{\|\epsilon\|_p \leq \rho} L_S(w + \epsilon)$$

- ▶ The new objective leads to a min-max problem:

- **Inner Maximization:** Find the "worst-case" loss within a small neighborhood of radius ρ around the current weights w . This tells us how sharp the current region is.
- **Outer Minimization:** Update the weights w to minimize that worst-case loss.

- ▶ One can interpret it as the worst-case optimization.

SAM: STEP 1 (ASCENT)

- ▶ **Step 1:** The first step is to solve the inner maximization problem:

$$\hat{\epsilon}(w) = \arg \max_{\|\epsilon\|_p \leq \rho} L_S(w + \epsilon)$$

- Finding the exact solution for this is infeasible.
- **Approximation:** Use a simple approximation based on a first-order Taylor expansion:

$$L_S(w + \epsilon) \approx L_S(w) + \epsilon^T \nabla_w L_S(w)$$

- To maximize this, the perturbation ϵ should point in the same direction as the gradient. The solution is a single gradient step of size ρ :

$$\hat{\epsilon}(w) = \rho \frac{\nabla_w L_S(w)}{\|\nabla_w L_S(w)\|_2}$$

SAM: STEP 2 (DESCENT)

- ▶ **Step 2:** After finding the "worst-case" perturbation $\hat{\epsilon}(w)$, we perform the main update.
- ▶ Instead of using the original gradient, SAM uses the gradient at the perturbed, "worst-case" point:

$$\nabla_w L_S^{\text{SAM}}(w) \approx \nabla_w L_S(w + \hat{\epsilon}(w))$$

- ▶ This leads to a simple two-step update process:
 1. **Ascent Step:** Calculate $\hat{\epsilon}(w)$ by taking a gradient step *uphill*.
 2. **Descent Step:** Calculate the gradient at this new point, $w + \hat{\epsilon}(w)$, and use it to update the weights w *downhill*.
- ▶ This means each SAM update requires two gradient computations, making it about twice as slow as a standard optimizer.

SAM: THEORETICAL JUSTIFICATION

- ▶ Why does this min-max objective work? This can be justified with a **PAC-Bayes generalization bound**.

Theorem (stated informally) 1. *For any $\rho > 0$, with high probability over training set \mathcal{S} generated from distribution \mathcal{D} ,*

$$L_{\mathcal{D}}(\mathbf{w}) \leq \max_{\|\epsilon\|_2 \leq \rho} L_{\mathcal{S}}(\mathbf{w} + \epsilon) + h(\|\mathbf{w}\|_2^2 / \rho^2),$$

where $h : \mathbb{R}_+ \rightarrow \mathbb{R}_+$ is a strictly increasing function (under some technical conditions on $L_{\mathcal{D}}(\mathbf{w})$).

- ▶ **Statement:** With high probability, the true test loss ($L_{\mathcal{D}}$) is bounded by the SAM loss on the training set ($L_{\mathcal{S}}$), plus a complexity term:
 - $L_{\mathcal{D}}(\mathbf{w})$: The true test loss.
 - The bound consists of two parts.
 - ▶ The first term, $\max_{\|\epsilon\|_2 \leq \rho} L_{\mathcal{S}}(\mathbf{w} + \epsilon)$, is **exactly the SAM loss objective**.
 - ▶ The second term is a complexity term that depends on the norm of the weights.
- ▶ **Implication:** By directly minimizing the SAM loss, the SAM algorithm is directly minimizing the main part of this generalization bound. This provides a theoretical connection between the SAM objective and improved generalization performance.

SAM: VISUALIZATION

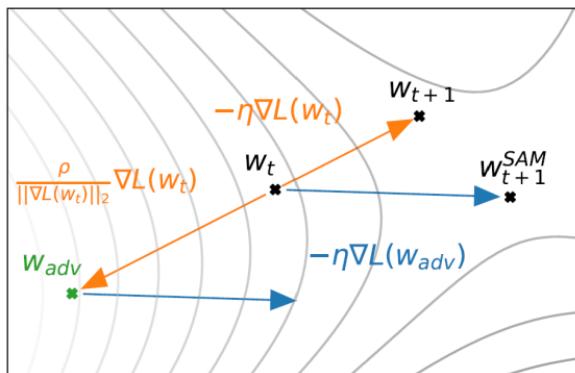
- ▶ This two-step process changes how the optimizer moves through the loss landscape.

SAM:

SGD:

- ▶ Directly follows the steepest descent w_{t+1} .
- ▶ Can easily fall into the nearest minimum, even if it is sharp.

- ▶ First, it "looks ahead" by stepping uphill to find a high point w_{adv} .
- ▶ Then, it pushes that high point down.
- ▶ This leads the optimizer away from sharp regions and towards wide, flat valleys w_{t+1}^{SAM} .

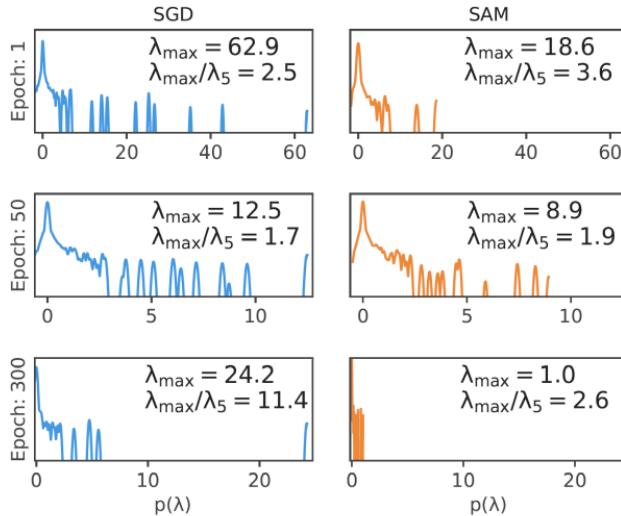


SAM: EMPIRICAL RESULTS

Model	Augmentation	CIFAR-10		CIFAR-100	
		SAM	SGD	SAM	SGD
WRN-28-10 (200 epochs)	Basic	2.7 \pm 0.1	3.5 \pm 0.1	16.5 \pm 0.2	18.8 \pm 0.2
	Cutout	2.3 \pm 0.1	2.6 \pm 0.1	14.9 \pm 0.2	16.9 \pm 0.1
	AA	2.1 \pm <0.1	2.3 \pm 0.1	13.6 \pm 0.2	15.8 \pm 0.2
WRN-28-10 (1800 epochs)	Basic	2.4 \pm 0.1	3.5 \pm 0.1	16.3 \pm 0.2	19.1 \pm 0.1
	Cutout	2.1 \pm 0.1	2.7 \pm 0.1	14.0 \pm 0.1	17.4 \pm 0.1
	AA	1.6 \pm 0.1	2.2 \pm <0.1	12.8 \pm 0.2	16.1 \pm 0.2
Shake-Shake (26 2x96d)	Basic	2.3 \pm <0.1	2.7 \pm 0.1	15.1 \pm 0.1	17.0 \pm 0.1
	Cutout	2.0 \pm <0.1	2.3 \pm 0.1	14.2 \pm 0.2	15.7 \pm 0.2
	AA	1.6 \pm <0.1	1.9 \pm 0.1	12.8 \pm 0.1	14.1 \pm 0.2
PyramidNet	Basic	2.7 \pm 0.1	4.0 \pm 0.1	14.6 \pm 0.4	19.7 \pm 0.3
	Cutout	1.9 \pm 0.1	2.5 \pm 0.1	12.6 \pm 0.2	16.4 \pm 0.1
	AA	1.6 \pm 0.1	1.9 \pm 0.1	11.6 \pm 0.1	14.6 \pm 0.1
PyramidNet+ShakeDrop	Basic	2.1 \pm 0.1	2.5 \pm 0.1	13.3 \pm 0.2	14.5 \pm 0.1
	Cutout	1.6 \pm <0.1	1.9 \pm 0.1	11.3 \pm 0.1	11.8 \pm 0.2
	AA	1.4 \pm <0.1	1.6 \pm <0.1	10.3 \pm 0.1	10.6 \pm 0.1

- SAM showed consistent improvements in generalization across many different models and datasets.

SAM: EMPIRICAL RESULTS



- ▶ The SGD solution has many large eigenvalues, especially at the end of training ($\lambda_{\max} = 24.2$). This indicates a sharp minimum.
- ▶ The SAM solution has a much smaller maximum eigenvalue ($\lambda_{\max} = 1.0$). This provides strong evidence that SAM successfully finds a much flatter minimum.

SAM: A DEEPER LOOK INTO SAM'S MECHANISM

- ▶ Three Kinds of Sharpness (2023 ICLR, Wen et al.):

- 1. Worst-Direction Sharpness (L_ρ^{Max}):

$$\max_{\|\epsilon\| \leq \rho} f(w + \epsilon)$$

This is the "true" sharpness that SAM originally wanted to minimize, which is related to the **largest eigenvalue** (λ_{\max}) of the Hessian.

- 2. Ascent-Direction Sharpness (L_ρ^{Asc}):

$$f\left(w + \rho \frac{\nabla f(w)}{\|\nabla f(w)\|_2}\right)$$

This is the approximation SAM uses for computational efficiency, which is related to the **smallest eigenvalue** (λ_{\min}) of the Hessian.

- 3. Average-Direction Sharpness (L_ρ^{Avg}):

$$\mathbb{E}_{g \sim N(0, I)} \left[f\left(w + \rho \frac{g}{\|g\|_2}\right) \right]$$

This is the sharpness measure used in the PAC-Bayes theory. It measures the average loss over random directions. It is related to the **trace** of the Hessian.

SAM: A DEEPER LOOK INTO SAM'S MECHANISM

- ▶ These three sharpness measures are **not the same**. Minimizing one can lead to a different solution than minimizing another.
- ▶ The SAM algorithm uses an approximation (L_ρ^{Asc}) that, in theory, should cause it to minimize the *smallest eigenvalue* λ_{\min} .
- ▶ However, it turns out that when using **full-batch** gradients, the SAM algorithm *actually* minimizes the **largest eigenvalue** (λ_{\max}), which corresponds to the "true" worst-direction sharpness (L_ρ^{Max}).
- ▶ More specifically, as training progresses, the gradient of SAM gradually becomes aligned with the top eigendirection of the Hessian, causing the updates to move along the direction that directly reduces λ_{\max} .
- ▶ On the other hand, stochastic SAM is actually minimizing the trace of the Hessian.

SAM: A DEEPER LOOK INTO ITS DYNAMICS

- ▶ We now analyze the behavior of SAM based on the paper (2023 JMLR, "The dynamics of SAM; bouncing across ravines and drifting towards wide minima").

SAM: A DEEPER LOOK INTO ITS DYNAMICS

- ▶ Consider a simple quadratic loss function, $f(w) = \frac{1}{2}w^T H w$.
- ▶ For a standard Gradient Descent (GD) optimizer, the updates would simply converge to the minimum at $w = 0$.
- ▶ However, SAM does not converge to $w = 0$. Instead, it converges to a repeating "bouncing" motion. The weights jump back and forth across the minimum.
- ▶ **Key Insight:** This oscillation always happens along the direction of the **largest curvature** (i.e., the direction of the top eigenvector of the Hessian, v_{max}). This "bouncing across the ravine" is the fundamental dynamic of SAM near a minimum.

SAM: A DEEPER LOOK INTO ITS DYNAMICS

- ▶ **Main question:** Why does “bouncing” lead to flatter minima?
- ▶ Near a minimum, the SAM update decomposes into
 1. **Large component (bouncing):** maintains a two-point oscillation along the top eigendirection v_1 (the sharpest-curvature direction).
 2. **Small component (drifting):** The iterate slowly drifts toward regions with smaller λ_{\max} (flatter minima).
- ▶ **Mechanism:** The oscillation probes the neighborhood along v_1 ; the orthogonal drift uses local variation of H in that direction (a “third-derivative” effect) to reduce λ_{\max} .
- ▶ **Net effect:** Bounce for exploration (sharp direction) + drift for λ_{\max} reduction \Rightarrow convergence toward wide minima.

SAM: A DEEPER LOOK INTO ITS DYNAMICS

